

# Project Machine Learning

Diego Valdes

March 20, 2019

```
# 005_project_ml.R
# @version: 1
# @author: Diego Valdes
# @date: Feb 14, 2019
# IST 687
# Clean openpowerlifting.csv - https://www.kaggle.com/open-powerlifting/powerlifting-database
```

```
library(randomForest)
```

```
## Warning: package 'randomForest' was built under R version 3.5.2
```

```
## randomForest 4.6-14
```

```
## Type rfNews() to see new features/changes/bug fixes.
```

```
library(neuralnet)
```

```
## Warning: package 'neuralnet' was built under R version 3.5.2
```

```
rm(list=ls()) # clear work space
```

```
#dev.off(dev.list()["RStudioGD"]) # clear plots
```

```
#####
#                                     Prepare Data
#####
```

```
setwd("C:/Users/dvjr2/Google Drive/Documents/Syracuse/IST_687/Project/")
```

```
filePath = "clean_power_lifting.csv" # cleaned data - yay Katie!
```

```
# read data and take a look
```

```
powerLiftingOG <- read.csv(file = filePath, header=TRUE, sep=",", stringsAsFactors = FALSE)
str(powerLiftingOG)
```

```
## 'data.frame':   316076 obs. of  13 variables:
## $ X             : int  1 2 4 5 6 7 9 10 11 12 ...
## $ MeetID        : int  0 0 0 0 0 0 0 0 0 0 ...
## $ Name          : chr   "Angie Belk Terry" "Dawn Bogart" "Dawn Bogart" "Destiny Dula" ...
## $ Sex           : chr   "F" "F" "F" "F" ...
## $ Equipment     : chr   "Wraps" "Single-ply" "Raw" "Raw" ...
## $ Age           : int  47 42 42 18 28 60 52 52 24 56 ...
## $ BodyweightKg  : num  59.6 58.5 58.5 63.7 62.4 ...
## $ WeightClassKg : chr   "60" "60" "60" "67.5" ...
## $ BestSquatKg   : num  47.6 142.9 NA NA 170.1 ...
## $ BestBenchKg   : num  20.4 95.2 95.2 31.8 77.1 ...
## $ BestDeadliftKg: num  70.3 163.3 NA 90.7 145.2 ...
## $ TotalKg       : num  138.3 401.4 95.2 122.5 392.4 ...
## $ Place         : int  1 1 1 1 1 1 1 1 1 1 ...
```

```
summary(powerLiftingOG)
```

```
##           X           MeetID           Name           Sex
## Min.      :    1   Min.      :    0   Length:316076   Length:316076
## 1st Qu.: 90561   1st Qu.:2516   Class :character   Class :character
## Median :179024   Median :5801   Mode  :character   Mode  :character
## Mean    :183718   Mean    :4962
## 3rd Qu.:274312   3rd Qu.:7018
## Max.    :386414   Max.    :8481
##
## Equipment           Age           BodyweightKg   WeightClassKg
## Length:316076   Min.      : 5.0   Min.      : 17.24   Length:316076
## Class :character   1st Qu.:22.0   1st Qu.: 70.00   Class :character
## Mode  :character   Median :28.0   Median : 82.80   Mode  :character
##                               Mean    :30.9   Mean    : 86.63
##                               3rd Qu.:36.0   3rd Qu.:100.00
##                               Max.    :95.0   Max.    :242.40
##                               NA's    :204920   NA's    :823
## BestSquatKg   BestBenchKg   BestDeadliftKg   TotalKg
## Min.      :-175.0   Min.      :-167.50   Min.      : 2.27   Min.      : 11.0
## 1st Qu.: 127.5   1st Qu.: 79.38   1st Qu.:149.69   1st Qu.: 287.5
## Median : 174.6   Median : 115.00   Median :195.00   Median : 435.4
## Mean    : 177.9   Mean    : 119.00   Mean    :195.57   Mean    : 433.9
## 3rd Qu.: 217.7   3rd Qu.: 150.00   3rd Qu.:237.50   3rd Qu.: 570.0
## Max.    : 573.8   Max.    : 455.86   Max.    :460.40   Max.    :1365.3
## NA's    :55095   NA's    :8299   NA's    :39720
## Place
## Min.      : 1.000
## 1st Qu.: 1.000
## Median : 1.000
## Mean    : 2.951
## 3rd Qu.: 3.000
## Max.    :77.000
##
```

```
powerLiftingOG$Sex = as.factor(powerLiftingOG$Sex) # Make Sex a factor for predicting
```

```
powerLifting = powerLiftingOG[ , c(-1,-2,-3, -5,-6, -8, -12, -13)] # get rid of unwanted cols
powerLifting = na.omit(powerLifting) # clear NAs to run model
```

```
powerLifting = powerLifting[sample(nrow(powerLifting)), ] # randomize
trngData = powerLifting[1:ceiling(nrow(powerLifting)*.7) , ] # trng data 70% of original data
testData = powerLifting[(ceiling(nrow(powerLifting)*.7)+1):nrow(powerLifting), ] # test data
```

```
#####
#                               Random Forest
#####
```

```
modelRF = randomForest(trngData[ , -1], trngData[ , 1]) # independent, dependent
summary(modelRF)
```

```
##           Length Class Mode
## call              3 -none- call
## type              1 -none- character
## predicted        181170 factor numeric
## err.rate          1500 -none- numeric
## confusion          6 -none- numeric
```

```

## votes          362340 matrix numeric
## oob.times      181170 -none- numeric
## classes        2 -none- character
## importance     4 -none- numeric
## importanceSD    0 -none- NULL
## localImportance 0 -none- NULL
## proximity      0 -none- NULL
## ntree          1 -none- numeric
## mtry           1 -none- numeric
## forest         14 -none- list
## y              181170 factor numeric
## test           0 -none- NULL
## inbag           0 -none- NULL

round(importance(modelRF), 2) # importance of variables, high is better

##              MeanDecreaseGini
## BodyweightKg      9521.10
## BestSquatKg       10644.46
## BestBenchKg       28598.59
## BestDeadliftKg    15228.76

# make predictions
results = predict(modelRF, testData[, -1])

# look at results w/ confusion matrix
table(results)

## results
##      F      M
## 17789 59854

table(testData$Sex)

##
##      F      M
## 18334 59309

matrix = table(results, testData$Sex)
matrix

##
## results      F      M
##      F 15103  2686
##      M  3231 56623

# male vs female accuracy
print(paste('Female Predict Accuracy:', matrix[1,1]/sum(testData$Sex == "F")))

## [1] "Female Predict Accuracy: 0.823770044725646"

print(paste('Male Predict Accuracy:', matrix[2,2]/sum(testData$Sex == "M")))

## [1] "Male Predict Accuracy: 0.954711763813249"

# Overall accuracy
print(paste('Overall Accuracy: ', (matrix[1,1] + matrix[2,2])/length(testData$Sex )))

## [1] "Overall Accuracy: 0.923792228533158"

```

```
#####
#                               Linear Model
#####

# change F/M to 1/0, LM seems to work better this way
trngData$Sex = ifelse(trngData$Sex == 'F', 1, 0)
testData$Sex = ifelse(testData$Sex == 'F', 1, 0)

# Lm with same variables used in randomForest()
modelGLM = glm(Sex ~ BodyweightKg + BestSquatKg + BestBenchKg + BestDeadliftKg, data = trngData, family = "binomial")
summary(modelGLM) # all variables are significant

##
## Call:
## glm(formula = Sex ~ BodyweightKg + BestSquatKg + BestBenchKg +
##      BestDeadliftKg, family = "binomial", data = trngData)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -3.3538  -0.3689  -0.1137  -0.0047   4.9663
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)    5.9136238  0.0433610  136.381 < 2e-16 ***
## BodyweightKg   -0.0024241  0.0004953   -4.894 9.88e-07 ***
## BestSquatKg     0.0076637  0.0004247   18.043 < 2e-16 ***
## BestBenchKg    -0.0686702  0.0006539 -105.016 < 2e-16 ***
## BestDeadliftKg -0.0113614  0.0004353  -26.102 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 197497  on 181169  degrees of freedom
## Residual deviance:  99006  on 181165  degrees of freedom
## AIC: 99016
##
## Number of Fisher Scoring iterations: 7

accuracy = .5 # accuracy variable, if prediction is above, assign prediction value
result = predict(modelGLM, testData[, 2:5], type = "response") # we want the percentage of females
result = ifelse(result > accuracy, 1, 0) # assign male or female prediction based on accuracy
Error = mean(result != testData[, 1]) # getting wrong values
print(paste('Overall Accuracy:', 1 - Error)) # total accuracy

## [1] "Overall Accuracy: 0.883917416895277"

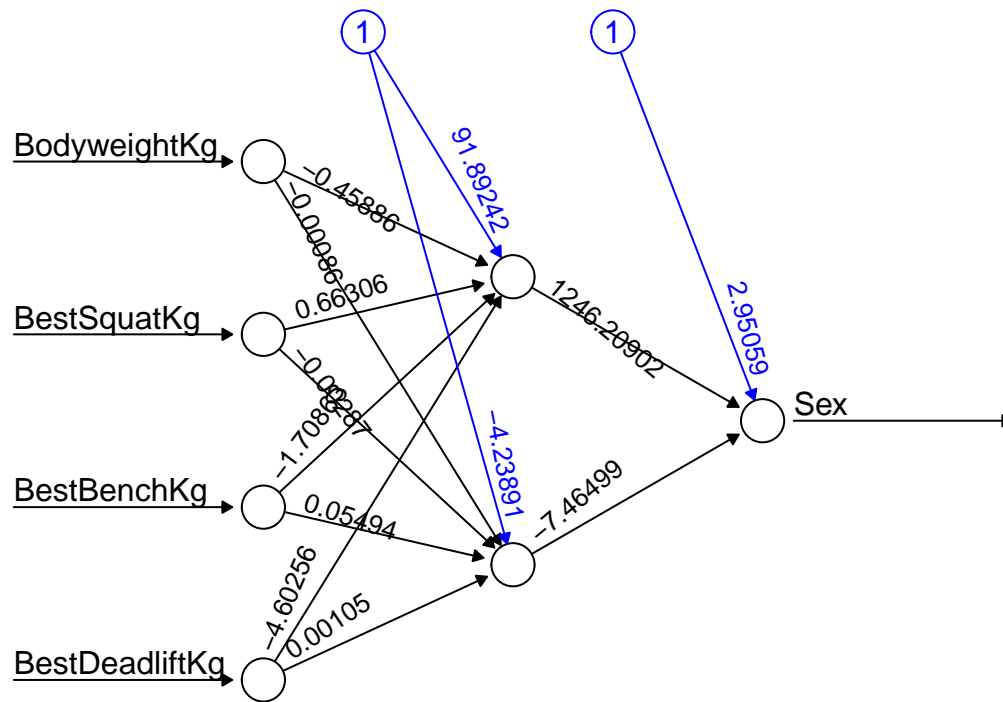
# build a neural network based on linear model
modelNeural = neuralnet(Sex ~ BodyweightKg + BestSquatKg + BestBenchKg + BestDeadliftKg, data = trngData,
                        hidden = 2, lifesign = 'minimal', linear.output = FALSE, threshold = 0.1)

## hidden: 2 thresh: 0.1 rep: 1/1 steps:
## 25944 error: 7412.50755 time: 18.02 mins
```

```
summary(modelNeural)
```

```
##                Length Class      Mode
## call           7 -none-      call
## response       181170 -none-    numeric
## covariate       724680 -none-    numeric
## model.list      2 -none-      list
## err.fct         1 -none-      function
## act.fct         1 -none-      function
## linear.output   1 -none-      logical
## data            5 data.frame list
## exclude        0 -none-      NULL
## net.result      1 -none-      list
## weights         1 -none-      list
## generalized.weights 1 -none-    list
## startweights    1 -none-      list
## result.matrix   16 -none-     numeric
```

```
plot(modelNeural, rep = 'best') #plot it
```



Error: 7412.507545 Steps: 25944

```
modelNeural$result.matrix # see the matrix
```

```
##                [,1]
## error          7.412508e+03
## reached.threshold 8.854398e-02
## steps          2.594400e+04
## Intercept.to.1layhid1 9.189242e+01
```

```
## BodyweightKg.to.1layhid1 -4.588611e-01
## BestSquatKg.to.1layhid1 6.630582e-01
## BestBenchKg.to.1layhid1 -1.708625e+00
## BestDeadliftKg.to.1layhid1 -4.602560e+00
## Intercept.to.1layhid2 -4.238908e+00
## BodyweightKg.to.1layhid2 -8.584729e-04
## BestSquatKg.to.1layhid2 -2.868649e-03
## BestBenchKg.to.1layhid2 5.494070e-02
## BestDeadliftKg.to.1layhid2 1.052004e-03
## Intercept.to.Sex 2.950593e+00
## 1layhid1.to.Sex 1.246209e+03
## 1layhid2.to.Sex -7.464994e+00

# predict and accuracy
resultNeural = predict(modelNeural, testData[ , 2:5], type = "response")
resultNeural = ifelse(resultNeural > accuracy, 1, 0) # assign male or female prediction based on accuracy
Error = mean(resultNeural != testData[ , 1]) # getting wrong values
print(paste('Overall Accuracy:', 1 - Error)) # total accuracy

## [1] "Overall Accuracy: 0.885385675463339"
```